

Generative AI and the Temporary Upskilling of Knowledge Workers

Emma Wiles^{1*}, Lisa Kraye^{2*}, Mohamed Abbadi², Urvi Awasthi²,
Ryan Kennedy², Pamela Mishkin³, Daniel Sack²,
Francois Candelon²

^{1*}Boston University, 595 Commonwealth Ave, Boston, 02215, MA.

²Henderson Institute, BCG, 200 Pier Four Blvd, Boston, 02110, MA.

³Economic Impacts Research, OpenAI, 3180 18th Street, SF, 94110, CA.

Abstract

Upskilling often refers to the process by which workers acquire and expand their skills, enabling them to perform different types of work as market demands change. This paper demonstrates that generative artificial intelligence can act as an “exoskeleton,” enhancing workers’ capabilities when attempting new skills, but these gains depend on continued technology use. We run a randomized controlled trial on upskilling by providing consultants (N = 573) access and training in using ChatGPT to solve data science tasks outside their skill-sets, which cannot be independently solved by ChatGPT. Treated workers scored higher than control group on three tasks, though we can reject the hypothesis that they performed as well as expert data scientists. We find no evidence treated workers are better at answering technical questions without ChatGPT post-experiment, suggesting their demonstrated technical capabilities do not imply knowledge acquisition. [†]

1 Introduction

The rapid advances in generative artificial intelligence (GenAI) and its widespread deployment have sparked both excitement and concern about its potential impact on the workforce. These models’ increasing capabilities in both automating and augmenting aspects of knowledge work raise questions about the impacts of technologies

*“Exoskeletons are wearable structures that support and assist movement, or augment the capabilities of the human body.” [1]

[†]Corresponding authors: emma.b.wiles@gmail.com, kraye.lisa@bcg.com. The full details of the experiment, including the survey text, can be found in our [supplementary information](#). The details of our preregistration can be found [here](#).

on the labor market—how jobs can be redesigned and how workers can adapt to the changing demands of employers. A key strategy for mitigating the negative effects of automation on workers is to “upskill” them into new and more in-demand skills [2, 3]. While recent studies have explored the effects of GenAI on worker performance in tasks within their existing skill sets [4, 5], there remains a gap in understanding how GenAI can be used to help workers acquire the new skills necessary to adapt to the changing demands of the labor market.

In general, we define skills in line with the task-based framework in labor economics, [6–8]. Specifically, we follow the definition of skills in Acemoglu and Autor (2011)— “a worker’s endowment of capabilities for performing various tasks.” Upskilling has been defined as workers gaining new skills to help in their current job responsibility [9]. Workers who upskill could have the task composition of their current job change. For example, a consultant who becomes a proficient programmer might take on more technical work without transitioning to become a data scientist or software engineer. We differentiate between two components of upskilling where—

- Workers gain capabilities in completing tasks, with aid from new technologies, or
- Workers acquire deeper knowledge, enabling them to complete tasks independently of technological aids.

In this paper, we find that GenAI enabled knowledge workers are “upskilled” according to only the first component. As a result, we ask the reader to consider the analogy of GenAI as an exoskeleton, where exoskeletons are: “wearable structures that support and assist movement, or augment the capabilities of the human body.” [1]. Like an exoskeleton, GenAI enables nontechnical workers to perform complex tasks outside their existing skill sets, but once the technology is removed, the enhanced capabilities disappear.

In this paper, we run a randomized controlled trial (RCT) involving 986 consultants at the BCG, a global management consulting firm. All participants were over the age of 20 and had, at minimum, bachelors degrees. These consultants are knowledge workers who likely know the basics of tools like Microsoft Excel, but the majority report not having coding experience. We show that providing nontechnical professionals with access to and training in GenAI chat-based tools (here, we use ChatGPT) can significantly enhance their ability to accurately complete data science tasks for problems that GenAI can perform and explain well but cannot independently solve.

Participants are randomly assigned to either a treatment group or a control group. In the first stage, all participants take a pre-experiment survey consisting of questions on participants’ subjective coding skills, GenAI usage, and a series of questions for which they had to predict whether or not GPT-4 could correctly answer them. Next, the treatment receives access to and training in ChatGPT, while the control group receives a training on using Stack Overflow and other public resources commonly used by data scientists. These 20 minute trainings both involve videos and interactive practice designed to familiarize participants with effective strategies for using their assigned tool for solving technical problems. Each participant is randomly assigned two of three 90-minute tasks. The coding task, statistical understanding task, and prediction task are all developed by BCG data scientists in collaboration with the Economic Impacts Research Team at OpenAI. These tasks are specifically designed to

be challenging for ChatGPT to solve independently, requiring human input at multiple steps. They were also designed to be outside the expectations of the skill-set of a general consultant such that if this task were required in a real-world project, a data-scientist would be staffed to do the work. Treated workers have access to ChatGPT’s latest model, GPT-4, and are encouraged to use it while attempting to complete the tasks. In the final stage, participants take a post-experiment survey that includes a series of technical questions on similar topics as the tasks they worked on.

For the 487 participants who submitted output for these tasks, we measure the treatment and control group workers’ performance by comparing their output to a benchmark set by the performance of data scientists (who perform the tasks without access to GenAI) who regularly perform these types of tasks. This approach allows us to directly test how much nontechnical workers armed with GenAI chatbots as their guide can do relative to workers in data science roles. While the expectations of data scientists vary from firm to firm, the tasks we study are representative of work that many data scientists commonly perform. However, it should be noted that we do not assess the impact of GenAI on workers’ performance on their ability to complete end-to-end data science problems nor do we test more advanced techniques like deep learning.

Treated workers perform better on all three tasks than the control group. On scores normalized to 0, which is the average performance of data scientists, treated workers perform 49, 20, and 17 percentage points better than those in the control group on coding, statistics, and prediction problems, respectively. We find that treated workers who begin the experiment with some coding skills perform just as well as do the data scientists on coding and statistics tasks. On the prediction problem, treated workers perform better than do control group workers, but there remains a large gap between their performance and the benchmark set by the data scientists.

We find no evidence that this successful practice leaves treated workers with more lasting knowledge. Following the experiment, on technical questions where the workers are not allowed to use ChatGPT, we find no evidence that workers in the treatment group perform better than the control group. Furthermore, we find that compared with those in the control group, treated workers exhibit greater overconfidence in the current capabilities of ChatGPT. In fact, following the experiment, treated workers perform worse than control workers at guessing which types of problems ChatGPT can and cannot solve.

We contribute to the new literature on the effects of GenAI on worker productivity and performance. Previous studies have demonstrated that GenAI can enhance efficiency and quality *within* workers’ existing skill sets. In one study, customer service agents given access to GenAI suggestions are able to resolve customer complaints more efficiently, and with no loss in quality, than are those without such access [5]. In another example, while workers with access to GenAI can complete tasks within the range of the model’s ability faster and more accurately than can workers without such access, they perform worse on tasks outside the model’s abilities [4]. We extend this literature by showing that GenAI can also improve the productivity of workers on complex tasks *outside* of their skill sets, addressing a critical question about whether GenAI can support upskilling by enabling workers to perform tasks in new and unfamiliar domains.

Second, we contribute to the literature on job training and upskilling workers. Automation from AI is a serious concern among academics and policymakers, with upskilling workers cited as one of the primary strategies for workers to keep from being displaced [2]. There is evidence that the earnings premium for technology-intensive college subjects declines faster than that for more general subjects [8], suggesting benefits for technical workers to be able to learn new skills outside of the formal education system. Given the evidence that GenAI is a general-purpose technology [10] and an effective teacher [11], it may allow workers to flexibly solve new types of problems as they emerge. However, our findings suggest that achieving these benefits may require complementary interventions to if the goal is durable skills without the use of GenAI.

Finally, we contribute to the literature on the limitations of human-AI interactions. Previous studies have shown that people who are given access to AI often cannot judge the quality of AI outputs [12]. Employers who are given access to AI-written first drafts of job postings produce more generic job postings that are less likely to lead to a hire [15]. In another hiring setting, recruiters opt to take AI’s suggestions, even when such suggestions are not correct [13]. Our finding that workers with training in ChatGPT are overconfident in the model’s abilities complements these findings, and we show that exposure to ChatGPT worsens the ability of workers to predict the model boundaries. These results highlight the risks of overreliance on GenAI and underscore the need for careful consideration of their long-term role in workforce development.

2 Results

Our first set of outcomes is related to how the workers performed on the three data science tasks. We measured workers’ likelihood of finishing each task and how they performed in terms of correctness and their process of completing each task. 44 BCG data scientists complete each task, and their average grades serve as the benchmark, normalized to 0, for each task. The coding task involved 10 distinct steps, and grades were determined by percentage of correct steps taken. The statistical understanding task included multiple-choice questions with both correct and incorrect answers and open-ended questions. For the prediction task, participants’ output is a vector of probabilities that were compared to the benchmarks set by the data scientists’ output (Please see Section 4.1 for details on the tasks and how they were graded.).

Our second set of outcomes tested whether workers appeared to retain knowledge without ChatGPT. We measured this through questions in the postexperiment survey, in which neither group was allowed to use ChatGPT to answer.

Third, we tested whether experience using ChatGPT helped workers better gauge the bounds of its abilities. In the pre- and postexperiment surveys, we provided workers with a series of problems and asked them “How likely is GPT-4 to solve this problem correctly?” [14]. We hypothesized that after completing these tasks, workers in the treatment group would be better at forecasting which types of problems ChatGPT can reliably solve compared to those in the control group.

2.1 Treated workers performed better on data science tasks

Figure 1 shows that treated workers were able to perform all three data science problems better than were their counterparts in the control group. We normalize the performance for each task so that 0 is equivalent to the average performance of the 44 data scientists.

On the coding task, treated workers performed 49 percentage points better than the workers in the control group ($\beta = 0.49$, 95% CI = [0.42,0.56], $p = 0.00$). In the 90 minutes allotted for the coding task, no worker in the control group achieved a perfect score, while five workers in the treatment group achieved a perfect score. Because of this low level of variation, we make the primary outcome for the coding task the percentage of steps that the worker took correctly, i.e., how many of the correct steps he or she took in the allotted time. On average, control group workers were 63 percentage points away from the performance of data scientists, while treatment group workers were only 14 percentage points away.


On the statistical understanding problem, the treatment group performed 20 percentage points better than the control group ($\beta = 0.20$, 95% CI = [0.15,0.25], $p < 0.001$). The treated group performed on average only 12 percentage points lower than the benchmark, although we can reject with 95% confidence that they performed as well as the data scientists. In the control group, workers' performance was on average 32 percentage points lower than that of the data scientists.

On the prediction problem, compared with control group workers, treated workers obtain results 17 percentage points closer to those of data scientists ($\beta = 0.17$, 95% CI = [0.09,0.25], $p < 0.001$). We use the outputs provided by the data scientists as the benchmarks since there are multiple valid predictive models and no single ground truth. Each worker's score for the prediction problem is the minimum absolute error between his or her vector and the benchmark vectors. For ease of interpretation, we multiply this error term by -1 so that a score of 0 implies that workers submitted a correct vector and a score of -1 denotes a vector that is orthogonal to any baseline answer. The control group had an average score of -0.43. The treatment group performed better, with an average score of -0.26, but we can reject with 95% confidence that they performed as well as did the data scientists.

2.2 We cannot reject the hypothesis that workers with prior coding experience performed as well as the data scientists on two of the three tasks

Treated workers who reported having prior coding experience performed almost as well as the data scientists on two of the three tasks, as shown in Figure 2. In the control group, we see that prior coding ability is associated with better performance in coding and statistics tasks, while on the prediction task, control group workers of all coding backgrounds performed similarly. On the statistics problem, the treatment effects are largest for those with the least coding experience on the statistics and coding problems (There are examples in the literature that GenAI can have differential treatment effects by skill level, both suggesting a skill bias where the highest skilled benefit the most [16] and an inverse skill bias where the lowest skilled benefit the most

Fig. 1 Effect of AI treatment on workers’ ability to solve data science problems




Notes: This plot reports the effect of the treatment on the workers’ performance across the 3 tasks. The x-axis is the mean score for each treatment group on each set of problems, where 0 is the average for data scientists. The first outcome is the percentage of correct steps they take in answering the coding question. The second outcome is the sum of the workers’ scores on each statistics question, divided by the total number of possible points. The third outcome is the score they received on the prediction problem, which is -1 times mean absolute error. The error bars indicate the 95% CIs of the mean of the treatment group relative to the mean in the control group. All standard errors are Huber–White robust. The sample includes all experimental participants who submitted a portion of each task for grading. Sample sizes per task: coding (N=300; Treatment=167, Control=133), statistics (N=330; Treatment=175, Control=154), prediction (N=298; Treatment=162, Control=136). The tasks are described in Appendix Section A.2 and the full text is in Supplementary Information A.1. Grading for each problem can be found in Supplementary Information Section A.2. Regression details for comparing treatment and control can be found in Appendix Table B3.

[5, 17, 18]. In Appendix Tables B10 and B13 we show the results of the treatment interacted with workers prior coding ability and whether or not they have a degree in a quantitative field. We do not see compelling evidence of either a skill bias or an inverse skill bias.).

Among those with prior coding experience, or “competent coders,” the control group’s average score on the coding task was 56 percentage points away from the data scientists’ benchmark, while the treated workers scored a mere 5 percentage points less than the data scientists, or a treatment effect of 51 percentage points ($\beta = 0.51$, 95% CI = [0.38,0.64], $p < 0.001$). We find no evidence of a difference between the scores of the competent coders in the treatment group, and of the data scientists. The Bayes Factor is 0.25, suggesting evidence for the null hypothesis. On the statistics task, the control group’s competent coders performed 21 percentage points lower than the benchmark, while the treated workers scored only 5 percentage points lower, or a treatment effect of 16 percentage points ($\beta = 0.16$, 95% CI = [0.08,0.23], $p = <0.001$). Again, when comparing the competent coders in the treatment group to the data scientists, the difference in scores is not statistically significant and with a Bayes Factor of 0.37. On the prediction task, the treatment still increased the scores of the

competent coders by 27 percentage points ($\beta = 0.27$, 95% CI = [0.13,0.40], $p < 0.001$) but they remained 20 percentage points away from the data scientists' benchmark.

Fig. 2 Effect of AI treatment on workers' ability to solve data science problems according to their prior coding experience




Notes: This plot reports the effect of the treatment on the workers' scores on each data science task according to their self-reported prior experience with coding. The x-axis is the mean score for each treatment group on each set of problems, where 0 is the benchmark set by the data scientists and -1 means that the worker received no points on the task. The first outcome is the percentage of correct steps that they take in answering the coding question. The second outcome is the sum of the workers' scores on each statistics question, divided by the total number of possible points. The third outcome is the score they received on the prediction problem, which is -1 times their mean absolute error. Huber-White robust standard errors are based on the difference between the means of the treatment and control groups, and we plot two-sided 95% CIs around each estimate. Because the data scientists are all at least competent coders, we cannot compare the performance of each worker in the treatment group to that of the data scientists interacted with coding experience. The sample includes all experimental participants who submitted a portion of each task for grading. N per cell: Never coded (Control: N = 67; Treatment: N = 79), Basic coding (Control: N = 82; Treatment: N = 90), Competent coder (Control: N = 76; Treatment: N = 92). The tasks are described in Appendix Section A.2 and the full text is in Supplementary Information A.1. Grading for each problem can be found in Supplementary Information A.2. Regression details can be found in Table B10.

2.3 No evidence of an impact on workers' ability to answer technical problems without the help of ChatGPT

Despite evidence that the use of AI improved workers' performance on the data science problems, after ChatGPT is taken away, the treated workers are no more likely to be able to correctly answer questions about probabilities, machine learning, or coding. In the postexperiment survey, workers are asked five questions on topics related to

227 their tasks. Workers in both groups are instructed not to use ChatGPT to answer
 228 these questions. The results in Figure 3 show no evidence that the treated group
 229 performs better than does the control group on these questions. A joint test of the
 230 treatment effects across all five questions confirms that the treatment does not have
 231 a statistically significant impact ($F(5, 198) = 0.83, p = 0.531$). Individually, the point
 232 estimates for the treatment effects range from -7.7 percentage points to 4.8 percentage
 233 points. A Bayes Factor analysis for each of the individual questions provides moderate
 234 evidence supporting the null hypothesis of no difference between the treatment and
 235 control group, with Bayes Factors of 0.10, 0.18, 0.50, 0.13, and 0.62, respectively. (To
 236 ensure robustness, we examine the heterogeneous treatment effects based on whether
 237 the worker performed a task directly related to the type of question. We show the
 238 results in Table B14 of the Appendix, which provide no evidence of a differential
 239 treatment effect.)

Fig. 3 Effects of AI treatment on post-experiment data science knowledge without the use of GenAI



Notes: This analysis looks at the effect of treatment on workers' ability to correctly answer data science questions after the conclusion of the experiment. The x-axis is the mean probability of getting the correct answer for each treatment group. The y-axis has the text of each question, with the format of the answer. We report Huber-White (HC0) robust standard errors and two-sided 95% confidence intervals. The text of the questions can be found in Supplementary Information B.2. Sample sizes per question: Q1 (filtering rows): N=323; Q2 (coin toss probability): N=270; Q3 (distance-based algorithms): N=270; Q4 (missing data techniques): N=151; Q5 (prediction category): N=270. Regression details can be found in Appendix Table B7.

2.4 Treated workers exhibit overconfidence in AI’s current capabilities

Workers in the treatment group perform worse at estimating whether something is within ChatGPT’s capabilities after the experiment. This analysis was not pre-registered and should be considered exploratory. Before and after the experiment, we pose seven questions (or prompts) and ask the workers to report their opinions on the likelihood of ChatGPT correctly answering the questions [14]. ChatGPT is unable to answer 5 of the technical questions given in the pre-experiment survey and 4 of the questions given in the postexperiment survey.


Prior to the experiment, workers in each group have comparable levels of confidence in ChatGPT’s ability to correctly answer similar questions. After the experiment, however, workers in the treatment group become significantly more optimistic and more incorrect about ChatGPT’s capabilities, as shown in Figure 4. Workers in both groups are optimistic about ChatGPT’s capabilities—the base rate for each question is between 65 and 78%. For all four of the questions that ChatGPT cannot answer, treated workers are significantly more likely to report that they believe ChatGPT can correctly answer each question. Surprisingly, the only two questions where the difference between the groups is not statistically significant are on the questions that ChatGPT *can* correctly answer.

Given fears about overreliance on GenAI causing harm, we may be concerned that workers who are overconfident in ChatGPT perform worse or exhibit smaller treatment effects as they are less likely to effectively check the system’s work. In Appendix Table B12, we interact the treatment effects with the pre-experiment measure of overconfidence. We define overconfidence based on their answers to the pre-experiment survey, where we ask workers to estimate whether the ChatGPT can solve various problems. Most of the workers exhibit overconfidence. For the 5 problems that ChatGPT cannot solve, 80% of the workers believe that ChatGPT can solve at least 3 problems. Therefore, we label workers as overconfident if they believe that ChatGPT can solve at least 4 of the 5 problems that it cannot solve. We find that overconfident workers have slightly weaker treatment effects to their grades on the tasks than workers who are less overconfident; however, the former workers still exhibit large positive treatment effects.

3 Discussion

We run an RCT in the setting of a global management consulting firm to understand to what extent GenAI can be used to help nontechnical knowledge workers perform tasks outside their current skill set, in this case, in performing data science tasks. We find that workers given access to and training in ChatGPT significantly outperform those in the control group, with the largest effects seen on the coding task. Treated workers perform better on three different margins: treated workers are more likely to complete tasks, complete tasks in less time, and obtain higher scores conditional on completing tasks. In addition, workers perform nearly at the level of real data scientists on the coding and statistics tasks. Despite these gains in the treated workers technical

Fig. 4 Effect of AI treatment on workers' predictions about the capabilities of ChatGPT



Notes: This analysis looks at the effect of treatment on workers' confidence in AI's capabilities. The x-axis is the difference in workers' confidence out of 100 that ChatGPT can answer the question correctly. A summary of the text of each question is on the y axis, with whether or not ChatGPT can obtain the correct answer. A 95% CI is plotted around each treatment effect, based on robust standard errors. The text of the questions can be found in Supplementary Information B.2 [14]. Sample sizes per question: Q1: N=454; Q2: N=475; Q3: N=473; Q4: N=464; Q5: N=465; Q6: N=431; Q7: N=451. Regression details can be found in Appendix Table B8.

capabilities, we find no credible evidence of a difference in the levels of knowledge retention between the treated and control groups.

Moreover, we find that exposure to ChatGPT induces overconfidence in AI's abilities, with treated workers being more likely to believe that ChatGPT can solve problems that it actually cannot. This finding echoes those from other recent studies on human susceptibility to AI errors and overreliance on AI assistance [12, 15] and suggests that organizations think carefully about monitoring when workers are using GenAI to complete tasks outside their skill sets. Despite this finding, we do not see evidence that workers who are overconfident in GenAI's capabilities prior to the experiment perform any worse on the tasks than do other workers. While this suggests that workers' overconfidence does not cause any changes in their performance, in settings where the consequences of errors are more significant, organizations should implement quality control measures such as evaluating the output from GenAI-assisted workers.

These results suggest that demonstrating new capabilities is not equivalent to gaining new knowledge, but nevertheless highlight the potential for AI itself to help

workers adapt to the changing skill demands of the labor market. We provide empirical evidence of how AI-enabled upskilling can be used to help workers adapt to more technically demanding work. Such upskilling may not immediately enable workers to perform at the level of an expert on complex tasks, the evidence suggests that the effort required to reach expert level may be reduced. For example, we find no evidence of a difference between the scores of the competent coders in the treatment group, and of the data scientists.

There is a large stream of literature that shows that GenAI increases productivity for tasks within knowledge workers’ skill sets [4, 5, 19]. Here, we provide evidence that it can also be used to help workers perform new tasks outside their skill set. These results have implications for both workers and managers.

For knowledge workers, the results suggest that using GenAI can increase the range of work they can proficiently accomplish. For hiring managers, the results suggest that they consider a wider pool of applicants who are proficient in GenAI. Managing this likely necessitates a reorganization of teams. For instance, teams may benefit from including someone with a deep understanding of data science to oversee and evaluate the work produced by less technical workers.

Our results also point to some important limitations of GenAI use for work outside one’s skill set. While treated workers can complete the data science tasks with the aid of ChatGPT, they do not demonstrate any greater retention of data science knowledge. This finding suggests that there may be limits to the depth of genuine skill acquisition, at least in the short term. Prior research indicates that practice and repetition [20, 21] are crucial for skill acquisition. The education literature suggests that while intensive, focused learning can produce some immediate gains like higher test scores, these often fail to translate into long-term retention without reinforcement and practice [22, 23]. Although it is outside the scope of this study, the impacts on long-term learning are likely different after workers practice and are exposed long term to GenAI.

We extend the literature on GenAI’s productivity impacts to provide evidence that GenAI can also be used to widen the scope of work that workers can perform. However, there are some important limitations to these results. First, our experiment focuses on discrete technical tasks rather than complete data science roles. The tasks we study are representative of work that data scientists regularly perform, such as making predictions, cleaning and building data, and evaluating ML model outputs, but do not include more advanced techniques, such as deep learning, nor do they capture the full complexity of end-to-end data science projects. Second, our experimental design does not have the power or scope to assess whether workers can actually transition to new jobs or industries based on these new skills. While we demonstrate that workers can perform technical work with GenAI assistance, we cannot necessarily conclude that this translates into career transitions. Furthermore, the sample the experiment is run on, while not being technical workers, are highly educated knowledge workers. This poses a limitation to the generalizability of our results to less educated segments of the population.

Despite these limitations, our findings highlight the potential of GenAI as a tool for exoskeleton-like skill augmentation. Fully realizing this potential while mitigating the risks of overconfidence and overreliance require ongoing research to aid in responsible

342 implementation. Future research should address whether longer-term exposure to and
343 practice with GenAI can help workers gain not only new capabilities but also a deeper
344 learning of skills.

345 4 Methods

346 4.1 Experimental design

347 This research was done in compliance with all relevant ethical regulations. The project
348 was deemed exempt by the WCG IRB’s IRB Affairs Department under 45 CFR §
349 46.104(d)(3) and deemed not to be subject to a review by MIT’s IRB COUHES.

350 We report the results from a large RCT run on consultants of BCG, a global man-
351 agerial consulting firm, to test whether nontechnical knowledge workers can perform
352 data science work with the help of ChatGPT-4, [preregistered](#) on March 13, 2024, prior
353 to data collection. All analyses reported here adhere to the pre-registered plan, with
354 the exception of the results to workers ability to guess which questions ChatGPT could
355 correctly solve, which were exploratory. We report the results to all pre-registered
356 outcomes in either the main body of this paper or in Appendix B (One outcome we
357 pre-registered but do not have results for was a “coding effort score” based on number
358 of bugs and how many of those bugs were resolved. However, we came across technical
359 difficulties in calculating “the resolution” of a coding bug (it is not straightforward
360 to measure whether the bug was resolved or if the participant simply moved on),
361 and did not continue down this path.). The experiment took place in late March and
362 early April of 2024. In the recruitment phase, all current BCG consultants were sent
363 an email inviting them to participate in a study on upskilling and GenAI. The email
364 incentivized participation with the following message: ”1) Be on the cutting edge!
365 After the study is published you can share on social media; until then it is entirely
366 confidential – in and outside of BCG; 2) Get “office contribution” credit counting for
367 CDC; 3) Get the chance to be selected for a small-group and exclusive discussion with
368 OpenAI and receive some unique merchandise.” Every participant received 4 hours of
369 office contribution credit which marginally informs their end of year bonus by amounts
370 that vary by their office location. BCG offices track time spent on office participation
371 across activities such recruiting, learning and development, mentorship, etc in order to
372 ensure people are recognized for going above and beyond in their work expectations.

373 For those who registered, we obtained informed consent, and the participants were
374 then surveyed on their demographics, programming and ChatGPT skills, technology
375 openness, trust in ChatGPT, and learning orientation [24–26]. Simultaneously, BCG
376 data scientists were invited to participate in a similar exercise, where they would sim-
377 ply complete the tasks used in the experiment without using AI. The output from
378 the 44 data scientists who participated served as the benchmark for the “typical
379 performance of a data scientist.”

380 After the registration survey, workers were randomly assigned to either a treatment
381 group or a control group. Random assignment was stratified across gender, location,
382 role (i.e., associate or consultant), coding skills, college degree (i.e., bachelor’s, mas-
383 ter’s, or Ph.D.), and experience with ChatGPT for coding. The initial sample of 986
384 workers was split equally into treatment and control groups, of which 573 workers

began the survey. Data collection and analysis were not performed blind to the conditions of the experiments. Our analysis sample consisted of 487 participants who each completed two tasks. Our pre-analysis plan specified a minimum detectable effect (MDE) of 1/3 standard deviation of scores on the three tasks. The standard deviations of scores in the control group for the coding, statistics, and prediction tasks were 0.25, 0.23, and 0.40, respectively. A minimum detectable effect of 1/3 standard deviation corresponds to approximately 8.2, 7.6, and 13.4 percentage points for the scores of the respective tasks. Based on our power calculations, with a significance level of 0.01% and power of 90%, we determined that a sample size of approximately 264 workers per task was required to detect an effect size of 1/3 standard deviation. In our experiment, 300 workers completed the coding task, 330 completed the statistics task, and 298 completed the prediction task, suggesting that our study was well-powered.

The experiment contained four phases (Appendix Figure A1):

1. **A pre-experiment survey** consisting of questions on participants' subjective coding skills and GenAI usage and a series of questions for which they had to predict whether or not GPT-4 could correctly answer them.
2. **A 15-20 minute tailored training** to each experimental condition. The treatment group received training on ChatGPT prompting for technical problems, while the control group received training on leveraging Google and commonly used websites such as Stack Overflow and Khan Academy. Both trainings involved a combination of videos and interactive practice.
3. **Three data science tasks** designed to be more technical than participants' current roles. These problems were outside the expected skill set of knowledge workers and often representative of those tasks often undertaken by data scientists.
4. **A postexperiment survey** that included a series of technical questions on similar topics as the tasks that all participants had to answer without the help of ChatGPT.

4.2 Task descriptions

The first task was a coding task testing workers' ability to clean data and answer questions about data using Python. Workers were required to write and submit their Python code Google Collab. These workers were likely to know the basics of data cleaning using tools like Microsoft Excel and Alteryx for their roles as consultants. However, more than 60% reported either having never coded before or knowing only the basics of coding (It should be noted that the BCG general consultants are not expected to be able to write code or perform complex data science tasks, these tasks are instead performed by BCG data scientists. General consultants are hired and promoted based on skills in analytical thinking and problem solving - such as the ability to frame and structure complex problems and conduct quantitative analysis using common tools such as Excel and Tableau. The most complex analytical capability that general consultants are expected to have is the ability to use Alteryx - a no code / low code data analytics and visualization platform.).

The second task measured statistical understanding, the ability to fact-check GenAI’s statistical analysis recommendations, and the ability to interpret the output of machine learning models. In this task, workers were guided through a series of potential GenAI outputs to create a model that predicted whether a couple would take out a mortgage based on historical data. They were then asked a series of questions on how they would respond to various situations while solving machine learning problems, such determining the cause of poor model performance.

The third task was a prediction task, which required workers to create a predictive model using historical data on international football games to develop an investment strategy. Their goal was to assess the “predictability” of their model; i.e., they needed to assess how reliable their model was for investment decisions. This task involved testing participants’ ability to understand predictive modeling and correctly identify when and how to apply machine learning models.

The tasks were created by BCG data scientists in collaboration with the Economic Impacts Research Team at OpenAI. They were designed such that ChatGPT is a useful guide, but is unable to correctly answer any of the questions independently. For all three tasks, if the participant let ChatGPT answer the question on its own, then the answer would be incorrect. Each task was intended to take 90 minutes to complete, and thus, to avoid fatigue, we assigned each participant two of the three tasks randomly, with the task order being randomized.

4.3 Analysis sample and primary outcomes

For most of the experimental results, we included workers who submitted answers for both their first and second tasks. This sample included 487 workers across the treatment and control groups. In the first panel of Table 1, we show attrition from the registration survey to the analysis sample. While there was some attrition at each stage, we find no evidence that the magnitude of the attrition was significantly different between the treatment and control groups. We show that those who completed the experiment and the attritors did not differ between the treatment and control groups on observables, as shown in Appendix Table B2. In case of any differential attrition on unobservables, we report Lee Bounds [27] on our main results in Appendix Table B15 and find that all the main results retain their significance.

Our first set of outcomes is related to how the workers performed on the three data science tasks. We measured workers’ likelihood of finishing each task, how long it took them to complete each task, and how they performed in terms of correctness and their process of completing each task.

For the coding task, there was a conservatively defined correct answer, which received a “1,” while any other answer received a “0.” To obtain the correct answer, a worker must take 10 correct steps, and thus, we also measured the percentage of the correct steps taken by the worker. (We preregistered that we would look at both whether or not the answer was correct and the workers “approach” in solving the problem. We focus on the later outcome (what percentage of correct coding steps the worker took) because of low variation in the binary indicator for whether or not the worker reached the correct answer.) The statistical understanding task included multiple-choice questions with both correct and incorrect answers and open-ended

Table 1 Building the experimental sample

<i>Flow from initial allocation into analysis sample</i>				
	<i>Total (N)</i>	<i>Treatment (N)</i>	<i>Control (N)</i>	<i>P-value</i>
Total workers allocated	986	493	493	
↪ began survey	573	298	275	0.337
↪ completed first task	511	270	241	0.200
↪ completed second task	487	260	227	0.135
<i>Effect of treatment on submitting each task</i>				
↪ completed coding task	300	167	133	0.019
↪ completed statistics task	330	175	154	0.326
↪ completed predictions task	298	162	136	0.341

Notes: This table reports means of pretreatment covariates by treatment and control. The first panel describes the flow of the sample from the allocation to the sample we use for our main experimental analysis. The complete allocated sample is described in the first line, with each following line defined cumulatively. Each worker was assigned two tasks, and the following lines compare the number of workers who submitted any work for each of the tasks. Those who completed both tasks made up the main experimental sample. Reported p-values in Panel A are two-sided chi-square tests of equality of group counts at each stage. The second panel provides the sample size for each treatment group and each task and shows treatment effects on whether or not workers completed each task. For each task, p-values are two-sided tests for equality of completion rates between treatment and control using the started-task sample as the denominator (task starters are defined by non-missing task-order assignment in the full-sample derived dataset).

questions. For the prediction task, participants’ output is a vector of probabilities that were compared to the benchmarks set by the data scientists’ output (see Section A.3 for details). We used GPT-4 Turbo API to automatically grade the parts of the three tasks that did not have a binary right or wrong answer, for example, the free response portion of the statistics task. This automation was rigorously validated by a data scientist (Details of the grading process are given in overview in Appendix Section A.3.1 with the full details in the Supplementary Information A.2.).

Our second set of outcomes tested whether workers appeared to retain knowledge without ChatGPT. We measured this through questions in the postexperiment survey, in which neither group was allowed to use ChatGPT to answer. Here, we hoped to test whether any new abilities that workers gained while using ChatGPT in a time-pressured environment could be retained without the tool.

Third, we tested whether experience using ChatGPT helped workers better gauge the bounds of its abilities. In the pre- and postexperiment surveys, we provided workers with a series of problems and asked them “How likely is GPT-4 to solve this problem correctly?” [14]. We hypothesized that after completing these tasks, workers in the treatment group would be better at forecasting which types of problems ChatGPT can reliably solve compared to those in the control group. The outcomes related to forecasting which types of problems ChatGPT can solve were not pre-registered.

490 5 Data availability

491 The data that support the findings of this study are available from the correspond-
492 ing authors (emma.b.wiles@gmail.com, kraye.lisa@bcg.com). Data is available on
493 request for non-commercial, academic review only, provided the requester is a certified
494 academic and is willing to sign a data usage agreement.

495 6 Code availability

496 The code used for analysis is publicly available in this [repository](#).

497 7 Acknowledgements

498 The authors thank BCG and the BCG Henderson Institute for facilitating the exper-
499 iment and providing access to the participant pool. We are grateful to the BCG
500 data scientists who served as benchmark participants, and to the Economic Impacts
501 Research Team at OpenAI for collaboration on task design.

502 8 Author contributions

503 E. Wiles and L. Kraye are joint first authors and the corresponding authors.
504 M. Abbadi is second author. U. Awasthi, R. Kennedy, and P. Mishkin are joint
505 third authors. D. Sack and F. Candelon are joint supervising authors. All authors
506 participated in the design of the Pre-Analysis Plan.


507 9 Competing interests

508 L. Kraye, M. Abbadi, U. Awasthi, R. Kennedy, D. Sack, and F. Candelon are or
509 were employees of BCG. P. Mishkin is an employee of OpenAI. E. Wiles was a paid
510 consultant for BCG during the running of the experiment.

Appendix A Experimental design and task overview

A.1 Experimental design

Fig. A1 Overview of the experimental design



Notes: The registration survey is in Supplementary Information Section B.1. The pre-experiment survey and training information can be found in Supplementary Information Section B.2. The text of all three tasks is summarized in Appendix Section A.2 with exact survey in Supplementary Information Section A.1. The postexperiment survey can be found in Supplementary Information Section B.2.

514 A.2 Task description

515 For the purpose of this study, we identify three types of tasks to test nontechnical
516 workers’ ability to perform data science-related work. The three tasks are independent
517 of each other and test different skills, including the ability to 1) write Python code
518 to perform data cleaning; 2) identify and correct modeling errors, apply statistical
519 metrics, interpret statistical plots, and calculate probabilities; and 3) build and assess
520 a predictive model. The details of each task are as follows:

521 Coding task

522 The coding task is designed to test the ability of a worker to write functioning
523 Python code. Workers are asked to clean, merge, and filter two datasets. This type of
524 task is required in the early stages of any data science exercise. We provide partici-
525 pants with detailed guidelines and tips to help them accomplish the task (see Appendix
526 A.3 for details) to isolate the skill of writing functioning code from that of knowing
527 the appropriate data cleaning steps. Specifically, participants are prompted to write
528 Python code to read two data files that contain information about customers, order
529 numbers, and products. They are then provided clear data quality and cleaning guide-
530 lines and tips for handling missing, junk, and duplicate data and asked to identify
531 the five customers with the highest average by cost in a certain period. This exercise
532 involved the following steps:

- 533 • Reading the data/creating a data frame
- 534 • Splitting combined string features into separate columns
- 535 • Handling feature types for filtering
- 536 • Detecting junk values
- 537 • Removing null and junk values
- 538 • Joining data
- 539 • Calculating basic metrics

540 Before beginning the task, we provide participants with written and video instruc-
541 tions on how to download and install Colab, Google’s integrated development
542 environment (IDE), to ensure that it is functioning properly (see Supplementary Infor-
543 mation Section B for details). We choose Colab as the IDE because it is relatively
544 simple to set up for nontechnical workers and allows us to access and trace the code
545 for evaluation after task completion.

546 We divide the task into ten distinct steps, as described in the coding task rubric
547 (Supplementary Information Section A.2.2). For phase analysis, we combine steps 1-3
548 as the “data cleaning” phase, steps 4-5 as the “data correction” phase and steps 6-10
549 as the “data merging and filtering” phase.

550 Heterogeneous treatment effects on the coding task, the robustness of the results
551 and the details of the phase analysis can be found in Appendix Sections B.2.1, B.2.2
552 and B.2.3, respectively.

554 Statistical understanding task

555 In the statistical knowledge task, we assess participants’ ability to understand
556 and apply statistical knowledge. Workers are guided through a series of potential
557 statistical analysis recommendations and other erroneous GenAI outputs including 1)
558

an erroneous analysis plan for creating a model that predicts whether a couple will take out a mortgage based on historical data, 2) analyzing and interpreting residual plots, 3) understanding the application of statistical metrics such as empirical risk to graphical outputs and 4) calculating situational probability. The responses to these questions include a mix of open-ended and multiple-choice answers. The questions are designed to encourage engagement and interpretation rather than simply finding and copying answers. Participants are explicitly instructed not to copy and paste the questions or images into ChatGPT or Google.

The questions are framed around the following three scenarios:

1. Given a data sample that contains financial and demographic information about partners who purchase a home, we ask participants to identify the steps needed to explore and preprocess the data, split the data, train and evaluate the model and then interpret various residual plots.
2. Given a certain distribution of data points, participants are asked how to classify additional data points, identify empirical risks and name potential classifiers that can create certain boundaries.
3. Participants are given a scenario involving a delivery truck that went missing, where the manager has to decide which route to choose based on situational probability.

In addition to the overall average treatment effects, we address the following questions:

- the type of statistical knowledge they test (i.e., debugging a series of steps in a proposed model by ChatGPT, analyzing and interpreting plots, and calculating situational probability)
- whether the responses are open ended or multiple choice, and
- whether the questions are visual or text based.

Section A.2.3 in the Supplementary Information provides the grading rubric for the statistical knowledge task.

Heterogeneous treatment effects on the statistical knowledge task, the robustness of the results and the details of treatment effects based on the type of questions can be found in Sections B.2.1, B.2.2 and B.2.4, respectively.

Prediction task

The prediction task aims to test participants' ability to conduct quantitative analysis and build a predictive model based on historical data. We provide a dataset containing information about soccer matches, including the date, home-team name, away-team name, tournament type, whether the game is at a neutral location, and the number of goals scored by the home and away teams. Participants are asked to quantify the predictability of each match in the dataset. Finally, they are asked to identify the most surprising match based on their predictability metric.

The deliverables of the problem-solving task include the participants' textual description of the methodology they employ to solve the problem and a tabular format of their results for the quantified predictability of each match. Unlike the coding and statistical knowledge tasks, this task is more "open ended", requiring participants to

choose their method to build the model and identify an appropriate outcome to measure predictability. The rubric for grading the prediction task is detailed in Section A of the Appendix.

Heterogeneous treatment effects on the prediction task, the robustness of the results and the details on the phase analysis can be found in Sections B.2.1, B.2.2 and B.2.5, respectively.

A.3 Task grading methodology

Each task is graded with quantifiable measures of the correctness of the answers and approach, depending on the hypothesis. Each task is graded on both the correctness of the answer and the steps the participant uses to solve the problem. Below, we describe the main outcomes for the correctness of the answer and for the process scores.

Coding task

There is one distinct correct answer for the coding assignment. Correctness is a binary measure where 0 is incorrect and 1 is correct. Second, we compare the output from the workers and data scientists to a rubric we create with 10 steps, where each step is necessary to obtain the correct score. As with the statistics task, we grade this against the rubric (shown in Supplementary Information Section A.2). The rubric scores are weighted correctness scores such that the final score is determined by a weighted sum across all answers as follows:

$$\text{Total correctness} = \sum_{i=1}^n (\text{Correctness of answer}_i \times \text{Complexity weight}_i) \quad (\text{A1})$$

where n is the total number of distinct questions and the correctness of the answer, and the complexity weighting is defined as the level of complexity of the question. Similarly, this score is between $(0, 1)$, where 1 means that the worker takes all of the correct steps.

Statistical understanding task

Each question in the statistics task is graded against the rubric (shown in Supplementary Information Section A.2). The rubric scores are weighted correctness scores such that the final score is determined by a weighted sum across all answers as follows:

$$\text{Total correctness} = \sum_{i=1}^n (\text{Correctness of answer}_i \times \text{Complexity weight}_i) \quad (\text{A2})$$

where n is the total number of distinct questions and the correctness of the answer, and the complexity weighting is defined as the level of complexity of the question. The complexity weightings are determined by asking several lead data scientists with more than 5 years of experience to rank the complexity of each question and average it across their answers. This correctness measure is bounded by $(0, 1)$, where 1 is a perfect score.

639

640 **Prediction task**

641 The problem-solving task is designed to have numerous possible answers, some of
 642 which are better than others. We use the answers submitted by the data scientists as
 643 the baseline/benchmark to grade the workers' results. Specifically, participants submit
 644 a predictability score for each match. We normalize participants' predictability scores
 645 for each match, score_i , and calculate a loss score for the answers submitted by the
 646 workers compared to the data science benchmarks, DS score_i . For each participant,
 647 we create a loss score defined as follows:

$$\text{Loss Score} = \frac{1}{n} \sum_{i=0}^n |\text{score}_i - \text{DS score}_i| \quad (\text{A3})$$

648 where n is the number of soccer matches in the dataset.

649 The final score we give workers on the prediction problem is -1 times the loss score
 650 so that the score is between $(-1, 0)$, where 0 is a perfect score.

651

652 **Outcome selection**

Fig. A2 Outcome description and selection for the main analysis

Task	Coding		Statistical knowledge	Prediction	
Outcome measures	Correctness score	Process score	Correctness score	Mean absolute error	Process score
Description	0 or 1 depending on whether workers got all 5 of customer ID's with the highest average correct	Weighted number of steps taken correctly. There are 10-steps to get the correct answer. Each is weighed by complexity. This measure incorporates the correctness score.	Rubric for right/wrong answers created and verified by data scientists. Each score is weighted by complexity as determined by several lead data scientists.	Mean absolute difference between the predictability vectors submitted the workers to those submitted by any of the data scientists	Point based rubric which scores across factors such as choice of model (e.g. ML receive more points) and how they define predictability (e.g. Z-score receives the most points)
Selected for main analysis	✓		✓	✓	
Reasoning	We pre-registered the correctness score. However, only 5 workers in the treatment got the right answer and none in the control given time constraints, therefore we measure how much of the problem they complete in the timeframe		This score was pre-registered, rubrics were rigorously designed based on consensus of several lead data scientists with >5 years of experience	Mean absolute error selected as per pre-registration and verified by analyzing process scores, while there are many ways to solve this problem this demonstrates whether the worker was able to achieve a result comparable to a data scientist	

Copyright © 2024 by Boston Consulting Group. All rights reserved.

Notes: This figure illustrates the different outcomes for each task, those selected for our main analysis and the reasoning behind the selection.

653 A.3.1 Grading details

654 The validation process across all tasks was designed to ensure technical rigor
655 while addressing task-specific challenges. We implemented a rigorous validation pro-
656 cess (measuring accuracy and consistency with both human manual validation and
657 pre-defined metrics) designed to ensure both accuracy and consistency:

- 658 • **Randomized Batch Grading:** Participant responses were graded in random-
659 ized batches of five to reduce the risk of contextual bias. Statistical variability
660 within batches was closely monitored and informed improvements to grading
661 prompts.
- 662 • **Detailed Rubrics:** Grading was guided by comprehensive, task-specific rubrics
663 created and validated by Data Scientists and Lead Data Scientists (Data Scien-
664 tists with 4+ years of experience at BCG, likely more in the industry). These
665 rubrics outlined clear expectations for each task and were iteratively refined to
666 handle diverse approaches while maintaining fairness.
- 667 • **Preprocessing to Mitigate Bias:** Free-text responses were paraphrased and
668 summarized to prevent GPT-4 Turbo from favoring verbose or overly pol-
669 ished submissions. This preprocessing step ensured a level playing field for all
670 participants.
- 671 • **Human Benchmarks:** Between 25% and 100% of responses, depending on the
672 task, were manually reviewed and graded by Data Scientists to establish a ground
673 truth. These manual grades were compared to the automated scores to validate
674 the accuracy of the system and refine grading strategies. Prompt Optimization:
675 Grading prompts were refined iteratively to reduce inconsistencies and improve
676 alignment with ground truth scores. Metrics like mean squared error were used to
677 measure and address discrepancies between human and automated evaluations.
- 678 • **Multi-Pass Grading and Manual Review:** Each submission was graded mul-
679 tiple times, and consensus or averaged scores were used when the results were
680 consistent. Submissions with high variability across evaluations were flagged for
681 manual review.

682 Please refer to the detailed grading explanation in Supplementary Information A.2.1
683 for further explanation on the LLM grading strategies, and the grading validation.

684 A.4 Estimation of treatment effects

685 Across each of these metrics, we employ Equation (4) to estimate the average treatment
686 effects based on ordinary least squares regression, where y_i is the dependent variable
687 (e.g., representing a quantifiable measure of output quality in the coding task and
688 efficiency of code), and T_{GPT} is the ChatGPT treatment dummy. Finally, X_i is a
689 set of covariates collected in the survey—office location, gender, tenure at BCG, and
690 whether the worker is a native English speaker.

$$y_i = \beta_0 + \beta_{\text{GPT}}T_{\text{GPT}} + \gamma X_i + \varepsilon_i \quad (\text{A4})$$

Appendix B Appendix tables and figures

B.1 Balance check

Table B1 Comparison of worker covariates by treatment assignment

<i>Flow from initial allocation into analysis sample</i>				
	<i>Total (N)</i>	<i>Treatment (N)</i>	<i>Control (N)</i>	<i>P-value</i>
Total workers allocated	986	493	493	
↪ began survey	573	298	275	0.336
↪ completed first task	511	270	241	0.199
↪ completed second task	487	260	227	0.135
<i>Pre-allocation attributes of final sample: N = 487</i>				
	Treatment mean: $\bar{X}_{TRT}$	Control mean: $\bar{X}_{CTL}$	P-value	
Female	0.369	0.37	0.985	
Bachelors Degree	0.238	0.291	0.192	
Masters Degree	0.677	0.604	0.092	
Doctorate	0.085	0.106	0.428	
Consultant	0.515	0.493	0.361	
Low Tenure	0.535	0.542	0.893	
Native English	0.471	0.423	0.343	
Office in Africa	0.019	0.018	0.896	
Office in Asia Pacific	0.135	0.115	0.505	
Office in Central or South America	0.019	0.004	0.14	
Office in Europe or Middle East	0.492	0.52	0.546	
Office in North America	0.335	0.344	0.835	
Code at Work	0.277	0.285	0.875	
Some Python Familiarity	0.335	0.331	0.922	
Familiar with ChatGPT for Coding	0.413	0.381	0.518	
Use ChatGPT Daily for Work	0.394	0.438	0.371	

Notes: This table reports the means and standard errors of various pretreatment covariates for the treatment and control groups. The first panel describes the flow of the sample from the allocation to the sample we use for our main experimental analysis. The complete allocated sample is described in the first line, with each following line defined cumulatively. Each worker is assigned two tasks, and the following lines compare the number of workers who submit any work for each of the tasks. Those who complete both tasks make up the main experimental sample. The second panel looks at the pre-allocation characteristics of the job seekers in the sample we use for our analysis, $N = 487$. We report the fraction of workers on their self-reported i) gender, ii) highest degree achieved, and iii) office location. The reported p values are for two-sided t-tests of the null hypothesis of no difference in means across groups.

Table B2 Comparing those who submit something for both tasks (primary analysis sample) to attritors

		Sample Mean	Attritor Mean	P-value
Control	Female	0.38	0.48	0.18
Treatment		0.37	0.43	0.42
Control	Office in Europe or Middle East	0.52	0.43	0.19
Treatment		0.50	0.57	0.38
Control	Native English speaker	0.50	0.49	0.86
Treatment		0.41	0.48	0.41
Control	New hire (<1 year)	0.51	0.40	0.11
Treatment		0.54	0.55	0.98
Control	Proficient coder or better	0.31	0.22	0.11
Treatment		0.32	0.20	0.09
Control	Never coded	0.30	0.37	0.35
Treatment		0.31	0.32	0.93
Control	At most 1 coding language	0.93	0.95	0.75
Treatment		0.96	1.00	0.00
Control	PhD	0.11	0.00	0.00
Treatment		0.09	0.05	0.26
Control	Uses ChatGPT daily for work	0.38	0.48	0.16
Treatment		0.45	0.50	0.52
Control	Familiar with prompt engineering	0.68	0.58	0.15
Treatment		0.67	0.68	0.91
Control	Never code for work	0.59	0.68	0.28
Treatment		0.61	0.69	0.42

Notes: This table reports the mean of various pre-experiment covariates amongst the primary analysis sample with those who attrit. The primary analysis sample is made up of workers who submitted anything to be graded for both of their two tasks. The sample here is made up of all workers who started the pre-experiment survey, $N = 573$. We run a two-sided Welch Two Sample t-test on each covariate for attritors and non-attritors, within each treatment group.

B.2 Treatment effects of AI

Table B3 Effects of AI to workers performance on data science tasks, relative to data scientists

	<i>Dependent variable:</i>		
	Coding Task Score	Stats Task Score	Prediction Task Score
	(1)	(2)	(3)
GenAI Treatment Assigned (Trt)	0.490*** (0.036)	0.201*** (0.026)	0.172*** (0.042)
Mean Y in Control Group	-0.63	-0.32	-0.43
P-value (Trt)	0.000	0.000	0.000
Observations	300	330	298
R ²	0.360	0.151	0.056

Notes: This table analyzes the effect of the treatment on the consultants ability to correctly answer questions. Each outcome is normalized relative to the performance of BCG data scientists. The first outcome is the percentage of correct steps they take in answering the coding question. The second outcome is the sum of the consultant's score on each statistics question, divided by the total number of possible points. The third outcome is the score they got on the prediction problem, which is -1 times their mean absolute error from the correct answer. Exact definition of grading for each problem can be found in Appendix Section A. All reported standard errors are two-sided 95% standard errors computed using Huber–White (HC0) robust variance. Significance stars and p-values indicate two-sided tests using Huber–White (HC0) robust variance: * p < 0.05, ** p < 0.01, *** p < 0.001.

Table B4 Effects of AI to whether or not they submit any answer on each task

	<i>Dependent variable:</i>		
	Stats Submitted	Prediction Submitted	Coding Submitted
	(1)	(2)	(3)
GenAI Treatment Assigned (Trt)	0.062 (0.033)	0.018 (0.035)	0.069 (0.040)
Mean Y in Control Group	0.85	0.87	0.78
P-value (Trt)	0.066	0.616	0.089
Observations	369	364	369
R ²	0.024	0.006	0.015

Notes: This table analyzes the effect of the treatment on the consultants submitting any answer to each question. Text of problems can be found in Appendix Section A. All regressions include controls for gender, location, native english status, and low tenure. Reported entries are coefficient estimates with two-sided 95% standard errors computed using Huber–White (HC0) robust variance. Significance stars and p-values indicate two-sided tests using Huber–White (HC0) robust variance: * p < 0.05, ** p < 0.01, *** p < 0.001.

Table B5 Effects of AI to whether or not they submit any answer on each task

	<i>Dependent variable:</i>	
	Task 1 Submitted	Task 2 Submitted
	(1)	(2)
GenAI Treatment Assigned (Trt)	0.026 (0.026)	0.043 (0.030)
Mean Y in Control Group	0.88	0.83
P-value (Trt)	0.315	0.151
Observations	573	573
R ²	0.014	0.016

Notes: This table analyzes the effect of the treatment on the consultants submitting any answer to each question. Text of problems can be found in Appendix Section A. All regressions include controls for gender, location, native english status, and low tenure. Reported entries are coefficient estimates with two-sided 95% standard errors computed using Huber–White (HC0) robust variance. Significance stars and p-values indicate two-sided tests using Huber–White (HC0) robust variance: * p < 0.05, ** p < 0.01, *** p < 0.001.

Table B6 Effects of AI on number of minutes to complete each task

	<i>Dependent variable:</i>		
	Mins on Stats	Mins on Prediction	Mins on Coding
	(1)	(2)	(3)
GenAI Treatment Assigned (Trt)	2.312 (2.407)	−14.450*** (2.946)	−8.884*** (2.521)
Mean Y in Control Group	63.61	68.36	78.01
P-value (Trt)	0.340	0.000	0.000
Observations	327	318	303
R ²	0.021	0.094	0.058

Notes: This table analyzes the effect of the treatment on the length of time it took for consultants to finish each task, conditional on completion. The outcome in Column (1) is the number of minutes they spent on the Statistics task. The outcome in Column (2) is the number of minutes spent on the Problem Solving and Prediction task. And the outcome in Column (3) is the number of minutes spent on the Coding task. Consultants were randomly assigned two of the three tasks, and given 90 minutes maximum to complete each. All regressions include controls for gender, location, native english status, and low tenure. All reported standard errors are two-sided 95% standard errors computed using Huber–White (HC0) robust variance. Significance stars and p-values indicate two-sided tests using Huber–White (HC0) robust variance: * p < 0.05, ** p < 0.01, *** p < 0.001.

Table B7 Effects of AI treatment on post-experiment data science knowledge (no AI use)

<i>Dependent variable:</i>					
	Data science or coding question				
	(1)	(2)	(3)	(4)	(5)
GenAI Treatment Assigned (Trt)	0.008 (0.038)	0.027 (0.027)	0.048 (0.027)	−0.006 (0.050)	−0.077 (0.040)
Mean Y in Control Group	0.35	0.91	0.73	0.80	0.78
P-value (Trt)	0.827	0.319	0.075	0.902	0.057
Observations	573	399	418	253	408
R ²	0.0001	0.003	0.008	0.0001	0.009

Notes: This table analyzes the effect of the treatment on the consultants ability to answer data science and coding questions, after the conclusion of the experiment. Full text of questions can be found in Supplementary Appendix B.2. Reported entries are coefficient estimates with two-sided 95% confidence intervals computed using Huber–White (HC0) robust variance. Significance stars and p-values indicate two-sided tests using Huber–White (HC0) robust variance: * p <0.05, ** p <0.01, *** p <0.001.

Table B8 Effects of AI treatment to post experiment questions about GPT-4's capabilities

<i>Dependent variable:</i>							
	“Can GPT-4 answer [this question] correctly?”						
	(1)	(2)	(3)	(4)	(5)	(6)	(7)
GenAI Treatment Assigned (Trt)	1.390 (2.002)	6.006* (2.351)	5.539* (2.428)	6.214** (2.287)	0.302 (2.348)	9.654*** (2.481)	5.340* (2.577)
Mean Y in Control Group	75.82	73.66	74.75	69.08	77.93	67.14	64.93
P-value (Trt)	0.488	0.012	0.025	0.007	0.898	0.000	0.040
Observations	454	475	473	464	465	431	451
R ²	0.042	0.024	0.027	0.025	0.005	0.086	0.013

Notes: This table analyzes the effect of the treatment on the consultants confidence in GPT-4's ability to get the right answer on various questions, after the conclusion of the experiment. For each question, the consultant gave a percentage confidence in GPT-4's ability to answer the question correctly. The question in Columns 1 and 4, and 5 GPT-4 usually get correct. Questions 2,3,6 and7, GPT-4 almost never get correct. Text of questions can be found in Appendix Section A. All regressions include controls for gender, location, native english status, and low tenure. All reported standard errors are two-sided 95% standard errors computed using Huber–White (HC0) robust variance. Significance stars and p-values indicate two-sided tests using Huber–White (HC0) robust variance: * p < 0.05, ** p < 0.01, *** p < 0.001.

Table B9 Effects of AI to workers performance on tasks, relative to data scientists

	<i>Dependent variable:</i>		
	Coding Task Score	Stats Task Score	Prediction Task Score
	(1)	(2)	(3)
Assigned Control	−0.629*** (0.081)	−0.323*** (0.044)	−0.432*** (0.034)
GenAI Treatment Assigned (Trt)	−0.140 (0.083)	−0.123** (0.044)	−0.260*** (0.024)
P-value (Assigned Control)	0.000	0.000	0.000
P-value (GenAI Treatment)	0.093	0.005	0.000
Observations	331	357	342
R ²	0.363	0.185	0.152

Notes: This table illustrates the performance of workers in the control and treatment arms compared to data scientists'. The first outcome is the percentage of correct steps they took on the coding problem, weighted by difficulty. The second outcome is the points the worker got on the statistics question, divided by the maximum points. The third outcome is their mean absolute error of the distance from their answer to the data scientist's answer. Exact definition of grading for each problem can be found in Appendix Section A. All reported standard errors are two-sided 95% standard errors computed using Huber–White (HC0) robust variance. Significance stars and p-values indicate two-sided tests using Huber–White (HC0) robust variance: * p < 0.05, ** p < 0.01, *** p < 0.001.

694 **B.2.1 Heterogeneous treatment effects of AI**

Table B10 Effects of AI to workers performance on data science tasks, relative to data scientists

	<i>Dependent variable:</i>		
	Coding Task Score	Stats Task Score	Prediction Task Score
	(1)	(2)	(3)
GenAI Treatment Assigned (Trt)	0.547*** (0.052)	0.311*** (0.048)	0.091 (0.087)
Coding basics	0.071 (0.043)	0.095* (0.042)	0.002 (0.082)
Competent coder	0.151** (0.049)	0.238*** (0.041)	-0.037 (0.092)
Coding basics x Trt	-0.131 (0.080)	-0.143* (0.063)	0.052 (0.109)
Competent coder x Trt	-0.038 (0.085)	-0.154* (0.062)	0.175 (0.112)
Constant	-0.708*** (0.026)	-0.444*** (0.032)	-0.422*** (0.064)
Mean Y in Control Group	-0.63	-0.32	-0.43
P-value (Trt)	0.000	0.000	0.296
P-value (Trt x Coding basics)	0.101	0.023	0.629
P-value (Trt x Competent coder)	0.652	0.014	0.118
Observations	300	330	298
R ²	0.387	0.243	0.070

Notes: This table analyzes the effect of the treatment on the consultants ability to correctly answer questions, by their pre-experiment coding knowledge. The omitted variable is “No prior coding experience.” Workers who report having basic coding ability are classified as “Coding basics” and those who report knowing how to code are classified as “competent coders.” Each outcome is normalized relative to the performance of BCG data scientists. The first outcome is the percentage of correct steps they take in answering the coding question. The second outcome is the sum of the consultant’s score on each statistics question, divided by the total number of possible points. The third outcome is the score they got on the prediction problem, which is 1- their mean absolute error. Exact definition of grading for each problem can be found in Appendix Section A. All standard errors are huber white robust. * p < 0.05, ** p < 0.01, *** p < 0.001.

Table B11 Effects of AI to workers performance on data science tasks, relative to data scientists

	<i>Dependent variable:</i>		
	Coding Task Score	Stats Task Score	Prediction Task Score
	(1)	(2)	(3)
GenAI Treatment Assigned (Trt)	0.503*** (0.041)	0.183*** (0.030)	0.128** (0.047)
Knowledge of GPT's strengths	0.017 (0.052)	0.003 (0.051)	-0.128 (0.085)
Knowledge of GPT's strengths x Trt	-0.062 (0.088)	0.080 (0.063)	0.195* (0.098)
Constant	-0.633*** (0.024)	-0.324*** (0.020)	-0.404*** (0.038)
P-value (Trt)	0.000	0.000	0.007
P-value (Trt x Knowledge of GPT's strengths)	0.483	0.204	0.048
Observations	300	330	298
R ²	0.361	0.160	0.069

Notes: This table analyzes the effect of the treatment on the consultants ability to correctly answer questions, by their pre-experiment coding knowledge. “Knowledge of GPT’s strengths” is 1 if the consultant got 4 out of 7, or more of questions correct on the pre-experiment survey asking about their guesses of whether or not GPT-4 can correctly answer a question. The omitted variable is consultants who got fewer tha 4 out of 7 correct. Each outcome is normalized relative to the performance of BCG data scientists. The first outcome is the percentage of correct steps they take in answering the coding question. The second outcome is the sum of the consultant’s score on each statistics question, divided by the total number of possible points. The third outcome is the score they got on the prediction problem, which is -1 times their mean absolute error. Exact definition of grading for each problem can be found in Appendix Section A. All standard errors are huber white robust. * p <0.05, ** p <0.01, *** p <0.001.

Table B12 Effects of AI to workers performance on data science tasks, relative to data scientists

	<i>Dependent variable:</i>		
	Coding Task Score	Stats Task Score	Prediction Task Score
	(1)	(2)	(3)
GenAI Treatment Assigned (Trt)	0.512*** (0.048)	0.250*** (0.038)	0.214*** (0.061)
Overconfident	0.026 (0.042)	0.041 (0.037)	0.088 (0.068)
Overconfident x Trt	-0.046 (0.072)	-0.094 (0.052)	-0.087 (0.083)
Constant	-0.642*** (0.029)	-0.345*** (0.028)	-0.475*** (0.050)
P-value (Trt)	0.000	0.000	0.001
P-value (Trt x Overconfident)	0.525	0.072	0.298
Observations	300	330	298
R ²	0.361	0.159	0.063

Notes: This table analyzes the effect of the treatment on the consultants ability to correctly answer questions, by their pre-experiment estimates of GPT-4's current capabilities. "Overconfident" is 1 if the consultant believed that GPT-4 could correctly answer at least four of the five questions which GPT-4 could not do. Each outcome is normalized relative to the performance of BCG data scientists. The first outcome is the percentage of correct steps they take in answering the coding question. The second outcome is the sum of the consultant's score on each statistics question, divided by the total number of possible points. The third outcome is the score they got on the prediction problem, which is -1 times their mean absolute error. Exact definition of grading for each problem can be found in Appendix Section A. All standard errors are huber white robust. * p < 0.05, ** p < 0.01, *** p < 0.001.

Table B13 Effects of AI to workers performance on data science tasks, relative to data scientists

	<i>Dependent variable:</i>		
	Coding Task Score	Stats Task Score	Prediction Task Score
	(1)	(2)	(3)
GenAI Treatment Assigned (Trt)	0.468*** (0.041)	0.184*** (0.030)	0.193*** (0.049)
Quant Degree	0.015 (0.052)	0.019 (0.042)	0.061 (0.075)
Quant Degree x Trt	0.090 (0.081)	0.089 (0.057)	-0.070 (0.093)
Constant	-0.634*** (0.023)	-0.329*** (0.021)	-0.450*** (0.041)
P-value (Trt)	0.000	0.000	0.000
P-value (Quant Degree x Trt)	0.269	0.120	0.452
Observations	300	330	298
R ²	0.367	0.167	0.059

Notes: This table analyzes the effect of the treatment on the consultants ability to correctly answer questions, by their pre-experiment estimates of GPT-4's current capabilities. "Quant Degree" is 1 if the worker reports having any formal degree in Statistics, Economics, Mathematics or Data Science. Each outcome is normalized relative to the performance of BCG data scientists. The first outcome is the percentage of correct steps they take in answering the coding question. The second outcome is the sum of the consultant's score on each statistics question, divided by the total number of possible points. The third outcome is the score they got on the prediction problem, which is -1 times their mean absolute error. Exact definition of grading for each problem can be found in Appendix Section A. All standard errors are huber white robust.

* p < 0.05, ** p < 0.01, *** p < 0.001.

Table B14 Effects of AI treatment to post experiment data science knowledge without use of AI

<i>Dependent variable:</i>					
2-6	Data science or coding question				
	(1)	(2)	(3)	(4)	(5)
GenAI Treatment Assigned (Trt)	0.0004 (0.052)	0.023 (0.032)	0.079* (0.034)	-0.067 (0.068)	-0.148** (0.049)
Assigned Coding Task	0.391*** (0.038)	0.914*** (0.023)	0.724*** (0.025)	0.817*** (0.049)	0.825*** (0.036)
P-value (Trt)	0.994	0.483	0.020	0.318	0.002
Observations	323	270	270	151	270
R ²	0.00000	0.002	0.020	0.007	0.033

Notes: This table analyzes the effect of the treatment on the consultants ability to answer data science and coding questions, after the conclusion of the experiment. The first problem is about coding, the second is about statistics, and the final three are about machine learning and prediction. For each problem the treatment is interacted with an indicator for whichever task gave the worker experience in that topic. Full text of questions can be found in Online Appendix B.2. Reported entries are coefficient estimates with **95% confidence intervals computed using Huber–White (HC0) robust variance**. Significance stars and p-values are based on the same HC0 robust variance. * p < 0.05, ** p < 0.01, *** p < 0.001.

Table B15 Lee Bounds on Treatment Effects for Main Results

	Treatment effect	Lee Lower Bound	Lee Upper Bound
Statistics Task Score	0.204***	0.197***	0.201***
P-value	< 0.001	< 0.001	< 0.001
Prediction Task Score	0.182***	0.164***	0.199**
P-value	< 0.001	< 0.001	0.005
Coding Task Score	0.491***	0.471***	0.555***
P-value	< 0.001	< 0.001	< 0.001

Notes: This table analyzes the effect of the treatment on the consultants ability to correctly answer questions. The first outcome is the sum of the consultant's score on each statistics question, divided by the total number of possible points. The second outcome is the percentage of correct steps they take in answering the coding question. The third outcome is the score they got on the prediction problem, which is their mean absolute error multiplied by -1 for ease of interpretation. Exact definition of grading can be found in Appendix Section A. All regressions include no controls. $N = 505$. Significance stars and p-values indicate two-sided tests using Huber–White (HC0) robust variance: * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

695 **B.2.2 Robustness of the results**

Table B16 Effects of AI to workers performance on coding task, relative to data scientists

	<i>Dependent variable:</i>				
	Coding Score				
	(1)	(2)	(3)	(4)	(5)
GenAI Treatment Assigned (Trt)	0.490*** (0.038)	0.491*** (0.038)	0.487*** (0.038)	0.559*** (0.069)	0.564*** (0.069)
Coding basics				0.102 (0.071)	0.074 (0.072)
Competent coder				0.179* (0.072)	0.129 (0.080)
Europe		0.053 (0.041)	0.072 (0.041)	0.075 (0.041)	0.088* (0.042)
Female		0.027 (0.040)	0.041 (0.041)	0.052 (0.040)	0.042 (0.041)
Low tenure		−0.042 (0.039)	−0.063 (0.040)	−0.056 (0.038)	−0.080 (0.041)
Native English speaker		0.051 (0.041)	0.023 (0.042)	0.055 (0.041)	0.025 (0.042)
Masters degree			−0.072 (0.046)		−0.074 (0.046)
PhD			−0.125 (0.072)		−0.117 (0.072)
Degree in quantitative field			0.048 (0.045)		0.040 (0.045)
STEM			0.063 (0.040)		0.045 (0.043)
Minimal experience with prediction			−0.102* (0.051)		−0.106* (0.051)
Some experience with prediction			0.046 (0.050)		0.0002 (0.056)
Data visualization experience			0.034 (0.049)		0.023 (0.049)
Coding basics x Trt				−0.162 (0.094)	−0.160 (0.095)
Competent coder x Trt				−0.042 (0.094)	−0.056 (0.094)
Constant	−0.629*** (0.028)	−0.668*** (0.050)	−0.653*** (0.082)	−0.781*** (0.071)	−0.690*** (0.093)
P-value (Trt)	0.000	0.000	0.000	0.000	0.000
Observations	300	300	300	300	300
R ²	0.360	0.369	0.409	0.403	0.425

Notes: This table analyzes the effect of the treatment on the worker's score on the coding task relative to the data scientists benchmark, where 0 is a perfect score. The first specification is the effect of treatment on their score. The second specification adds controls for gender, office location, native english status, and prior coding ability. The third specification is the same as column (2) but with additional controls for different measures of the workers prior technical experience. The fourth specification is the same as column (2), where the treatment is interacted with the workers prior coding experience. The fifth specification is the same as column (3), where the treatment is interacted with the workers prior coding experience. All reported standard errors are two-sided 95% standard errors computed using Huber–White (HC0) robust variance. Significance stars and p-values indicate two-sided tests using Huber–White (HC0) robust variance * p < 0.05, ** p < 0.01, *** p < 0.001.

Table B17 Effects of AI to workers performance on statistics task, relative to data scientists

	<i>Dependent variable:</i>				
	Statistics Score				
	(1)	(2)	(3)	(4)	(5)
GenAI Treatment Assigned (Trt)	0.201*** (0.026)	0.201*** (0.026)	0.206*** (0.026)	0.308*** (0.045)	0.309*** (0.045)
Coding basics				0.094* (0.047)	0.088 (0.048)
Competent coder				0.234*** (0.045)	0.210*** (0.050)
Europe		0.028 (0.029)	0.027 (0.029)	0.035 (0.028)	0.030 (0.029)
Female		−0.033 (0.027)	−0.010 (0.028)	−0.014 (0.026)	−0.011 (0.027)
Low tenure		0.036 (0.027)	0.051 (0.028)	0.032 (0.026)	0.042 (0.028)
Native English speaker		0.038 (0.029)	0.033 (0.030)	0.038 (0.028)	0.027 (0.030)
Masters degree			−0.027 (0.032)		−0.033 (0.032)
PhD			0.004 (0.052)		−0.015 (0.051)
Degree in quantitative field			0.041 (0.032)		0.032 (0.031)
STEM			0.0003 (0.028)		−0.024 (0.029)
Minimal experience with prediction			0.048 (0.036)		0.036 (0.035)
Some experience with prediction			0.121*** (0.034)		0.066 (0.037)
Data visualization experience			0.044 (0.035)		0.020 (0.035)
Coding basics x Trt				−0.137* (0.063)	−0.130* (0.063)
Competent coder x Trt				−0.147* (0.061)	−0.153* (0.062)
Constant	−0.323*** (0.019)	−0.362*** (0.036)	−0.461*** (0.061)	−0.491*** (0.046)	−0.499*** (0.065)
P-value (Trt)	0.000	0.000	0.000	0.000	0.000
Observations	330	330	330	330	330
R ²	0.151	0.165	0.224	0.254	0.273

Notes: This table analyzes the effect of the treatment on the worker's score on the statistics task relative to the data scientists benchmark, where 0 is a perfect score. The first specification is the effect of treatment on their score. The second specification adds controls for gender, office location, native english status, and prior coding ability. The third specification is the same as column (2) but with additional controls for different measures of the workers prior technical experience. The fourth specification is the same as column (2), where the treatment is interacted with the workers prior coding experience. The fifth specification is the same as column (3), where the treatment is interacted with the workers prior coding experience. All reported standard errors are two-sided 95% standard errors computed using Huber–White (HC0) robust variance. Significance stars and p-values indicate two-sided tests using Huber–White (HC0) robust variance * p < 0.05, ** p < 0.01, *** p < 0.001.

Table B18 Effects of AI to workers performance on a prediction task, relative to data scientists

	<i>Dependent variable:</i>				
	Mean Absolute Error				
	(1)	(2)	(3)	(4)	(5)
GenAI Treatment Assigned (Trt)	0.172*** (0.041)	0.182*** (0.041)	0.168*** (0.042)	0.103 (0.078)	0.086 (0.078)
Coding basics				0.014 (0.073)	0.043 (0.076)
Competent coder				-0.026 (0.080)	0.023 (0.096)
Europe		0.020 (0.046)	0.016 (0.046)	0.025 (0.046)	0.019 (0.047)
Female		0.023 (0.043)	0.011 (0.044)	0.028 (0.043)	0.008 (0.044)
Low tenure		-0.008 (0.042)	-0.026 (0.044)	-0.011 (0.042)	-0.028 (0.044)
Native English speaker		0.084 (0.047)	0.101* (0.049)	0.088 (0.047)	0.104* (0.049)
Masters degree			0.088 (0.050)		0.101* (0.050)
PhD			-0.084 (0.089)		-0.076 (0.089)
Degree in quantitative field			0.002 (0.048)		0.012 (0.049)
STEM			-0.016 (0.045)		-0.046 (0.048)
Minimal experience with prediction			-0.099 (0.054)		-0.094 (0.054)
Some experience with prediction			-0.037 (0.057)		-0.079 (0.065)
Data visualization experience			0.050 (0.052)		0.030 (0.053)
Coding basics x Trt				0.051 (0.101)	0.058 (0.101)
Competent coder x Trt				0.173 (0.107)	0.176 (0.107)
Constant	-0.432*** (0.030)	-0.491*** (0.055)	-0.516*** (0.087)	-0.494*** (0.077)	-0.515*** (0.098)
P-value (Trt)	0.000	0.000	0.000	0.240	0.311
Observations	298	298	298	298	298
R ²	0.056	0.068	0.101	0.083	0.120

Notes: This table analyzes the effect of the treatment on the worker's score on the prediction problem, -1 times their mean absolute error from the correct answer, where 0 is a perfect score. The first specification is the effect of treatment on their score. The second specification adds controls for gender, office location, native english status, and prior coding ability. The third specification is the same as column (2) but with additional controls for different measures of the workers prior technical experience. The fourth specification is the same as column (2), where the treatment is interacted with the workers prior coding experience. The fifth specification is the same as column (3), where the treatment is interacted with the workers prior coding experience. All reported standard errors are two-sided 95% standard errors computed using Huber-White (HC0) robust variance. Significance stars and p-values indicate two-sided tests using Huber-White (HC0) robust variance: * p < 0.05, ** p < 0.01, *** p < 0.001.

Table B19 Effect of AI treatment on workers' coding performance, relative to data scientists

	<i>Dependent variable:</i>		
	Data Cleaning	Data Correct	Merge & Filter
	(1)	(2)	(3)
GenAI Treatment Assigned (Trt)	0.204*** (0.032)	0.440*** (0.059)	0.738*** (0.061)
Mean Y in Control Group	-0.26	-0.64	-0.82
P-value (Trt)	0.000	0.000	0.000
Observations	300	300	300
R ²	0.129	0.172	0.343

Notes: This table reports the effect of the treatment on the consultants performance on the coding task. The y-axis is the worker's binary grade on each set of actions in python, where 0 is a perfect score and -1 is the lowest possible score. The first outcome is whether or not the worker correctly cleaned the data. The second outcome is whether or not they correctly handle nulls and duplicates. The third outcome is whether or not they correctly merge and then filter a dataset. The sample includes all experimental participants who submitted something for grading on the prediction task. Text of problems can be found in Appendix Section A. Exact definition of grading for each problem can be found in Appendix Section A. All regressions include controls for gender, location, native english status, and low tenure. Reported entries are coefficient estimates with two-sided 95% standard errors computed using Huber–White (HC0) robust variance. Significance stars and p-values indicate two-sided tests using Huber–White (HC0) robust variance: * p < 0.05, ** p < 0.01, *** p < 0.001.

B.2.4 Treatment effects on the statistics task – Details

Table B20 Effect of AI treatment on workers’ statistics performance, relative to data scientists

	<i>Dependent variable:</i>		
	Identifying Errors	Visual ML Questions	Probability Questions
	(1)	(2)	(3)
GenAI Treatment Assigned (Trt)	0.308*** (0.038)	0.068* (0.032)	0.302*** (0.050)
Mean Y in Control Group	-0.48	-0.22	-0.26
p-value (Trt)	0.000	0.033	0.000
Observations	330	330	330
R ²	0.183	0.035	0.109

Notes: This table reports the effect of the treatment on the consultants grades on each step of the statistics task. Each outcome is the worker’s grade on each set of problems, where 0 is a perfect score and -1 is the lowest possible score. The first outcome is the workers’ score on a problem where they apply statistical metrics and interpret their graphs. The second outcome is their score on questions where they look at data plotted and decide what types of classifier to use on it. The third outcome is their score on probability questions. The sample includes all experimental participants who submitted something for grading on the statistics task. Text of problems can be found in Appendix Section A. Exact definition of grading for each problem can be found in Appendix Section A. Reported entries are coefficient estimates with two-sided 95% standard errors computed using Huber–White (HC0) robust variance. Significance stars and p-values indicate two-sided tests using Huber–White (HC0) robust variance: * p < 0.05, ** p < 0.01, *** p < 0.001. All regressions include controls for gender, location, native english status, and low tenure.

B.2.5 Treatment effects on the prediction task – Details

Table B21 Effect of AI treatment on workers' prediction scores, relative to data scientists

	<i>Dependent variable:</i>	
	Mean Absolute Error	Process Score
	(1)	(2)
GenAI Treatment Assigned (Trt)	0.182*** (0.041)	0.129** (0.045)
Mean Y in Control Group	-0.43	-0.37
p-value (Trt)	0.000	0.004
Observations	298	281
R ²	0.068	0.036

Notes: This table reports the effect of the treatment on the consultants performance on the prediction task. Each outcome is the worker's grade on each set of problems, where 0 is a perfect score and -1 is the lowest possible score. The first outcome is the score they got on the prediction problem, which is -1 times their mean absolute error. The second outcome is the percentage of "correct" steps they took. The sample includes all experimental participants who submitted something for grading on the prediction task. Text of problems can be found in Appendix Section A. Exact definition of grading for each problem can be found in Appendix Section A. All regressions include controls for gender, location, native english status, and low tenure. Reported entries are coefficient estimates with two-sided 95% standard errors computed using Huber–White (HC0) robust variance. Significance stars and p-values indicate two-sided tests using Huber–White (HC0) robust variance: * p < 0.05, ** p < 0.01, *** p < 0.001. All regressions include controls for gender, location, native english status, and low tenure.

Table B22 Effects of GenAI on Self-Reported Capabilities and Professional Confidence, Before and After Experiment

	<i>Dependent variable:</i>		
	Δ in Beliefs in Data Analytics Abilities	Δ in Confidence in Data Science Skills	Δ in Confidence with GenAI
GenAI Treatment Assigned (Trt)	-0.305 (1.673)	0.651*** (0.114)	-0.087 (0.199)
Constant	-0.451 (1.170)	-0.343*** (0.086)	0.309* (0.143)
Observations	421	463	477
p-value (Trt)	0.855	0.000	0.662
R ²	0.000	0.067	0.001


Notes: This table analyzes the effect of treatment on three key outcomes. Column (1) measures the change in respondents' agreement with the statement "Generative AI enables my ability to execute data analytics tasks and reach desired outcomes" on a 1-5 scale from strongly disagree to strongly agree, as measured before and after the experiment. Column (2) measures the change in response to "How confident are you in your ability to contribute to data science projects?" on a 1-7 scale, as measured before and after the experiment. Column (3) measures the change in a composite measure of professional confidence with GenAI, combining responses to three statements: "Generative AI helps me feel more competent in my role", "Generative AI enables my ability to execute tasks and reach desired outcomes", and "Generative AI helps me be more confident that I will meet my project managers expectations". Changes are measured as post-experiment score on the survey question minus the pre-experiment score. Standard errors in parentheses. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Table B23 Effects of GenAI on Self-Reported Beliefs in ChatGPT’s Impact on Career Development, Before and After Experiment

	<i>Dependent variable:</i>	
	Δ in Beliefs about ChatGPT’s Usefulness for Problem Solving	Δ in Beliefs about ChatGPT’s Usefulness for Practicality & Effectiveness
GenAI Treatment Assigned (Trt)	-0.215* (0.086)	-0.201* (0.082)
Constant	0.242*** (0.062)	0.128** (0.056)
Observations	478	478
p-value (Trt)	0.013	0.015
R ²	0.013	0.012


Notes: This table analyzes the effect of treatment on beliefs about how sustained use of ChatGPT for data science would impact career development. Column (1) measures the change in agreement with the statement "Sustained use of ChatGPT for data science would have the potential to make me a better consultant in the 'Problem solving and insights' dimension". Column (2) measures the change in agreement with an analogous statement about the 'Practicality and Effectiveness' dimension. Both questions used a 5-point scale from strongly disagree (1) to strongly agree (5). Changes are measured as post-experiment score on the survey question minus the pre-experiment score. Standard errors in parentheses. * $p < 0.05$, ** $p < 0.01$, *** $p < 0.001$.

Fig. B3 Effect of AI treatment on workers' ability to solve data science problems




Notes: This plot reports the effect of the treatment on the workers' performance across the 3 tasks. The x-axis is the mean score for each treatment group on each set of problems, where 0 is the average for data scientists. The first outcome is the percentage of correct steps they take in answering the coding question. The second outcome is the sum of the workers' scores on each statistics question, divided by the total number of possible points. The third outcome is the score they received on the prediction problem, which is -1 times mean absolute error. The error bars indicate the 95% CIs of the mean of the treatment group relative to the mean in the control group. All standard errors are Huber–White robust. The sample includes all experimental participants who submitted a portion of each task for grading. The tasks are described in Appendix Section A.2 and the full text is in Supplementary Appendix A.1. Grading for each problem can be found in Supplementary Appendix A.2. Regression details for comparing treatment and control can be found in Appendix Table B3.

Fig. B4 Effect of AI treatment on workers' ability to solve data science problems according to their prior coding experience



Notes: This plot reports the effect of the treatment on the workers' scores on each data science task according to their self-reported prior experience with coding. The x-axis is the mean score for each treatment group on each set of problems, where 0 is the benchmark set by the data scientists and -1 means that the worker received no points on the task. The first outcome is the percentage of correct steps that they take in answering the coding question. The second outcome is the sum of the workers' scores on each statistics question, divided by the total number of possible points. The third outcome is the score they received on the prediction problem, which is -1 times their mean absolute error. Huber-White robust standard errors are based on the difference between the means of the treatment and control groups, and we plot two-sided 95% CIs around each estimate. Because the data scientists are all at least competent coders, we cannot compare the performance of each worker in the treatment group to that of the data scientists interacted with coding experience. The sample includes all experimental participants who submitted a portion of each task for grading. N per cell: Never coded (Control: N = 67; Treatment: N = 79), Basic coding (Control: N = 82; Treatment: N = 90), Competent coder (Control: N = 76; Treatment: N = 92). The tasks are described in Appendix Section A.2 and the full text is in Supplementary Appendix A.1. Grading for each problem can be found in Supplementary Appendix A.2. Regression details can be found in Table B10.

Fig. B5 Effects of AI treatment on post-experiment data science knowledge without the use of GenAI



Notes: This analysis looks at the effect of treatment on workers' ability to correctly answer data science questions after the conclusion of the experiment. The x-axis is the mean probability of getting the correct answer for each treatment group. The y-axis has the text of each question, with the format of the answer. We report Huber–White (HC0) robust standard errors and two-sided 95% confidence intervals. The text of the questions can be found in Supplementary Appendix B.2. Regression details can be found in Appendix Table B7.

References

- [1] **Iberdrola**, “What are exoskeletons and how can they help us overcome our human limitations?,” 2024. Available at: <https://www.iberdrola.com/innovation/what-are-exoskeletons>. Accessed: 2024-09-03.
- [2] **Djankov, Simeon and Federica Saliola**, “The changing nature of work,” *Journal of International Affairs*, 2018, 72 (1), 57–74.
- [3] **Capraro, Valerio, Austin Lentsch, Daron Acemoglu et al.**, “The impact of generative artificial intelligence on socioeconomic inequalities and policy making,” *PNAS Nexus*, 2024, 3 (6).
- [4] **Dell’Acqua, Fabrizio, McFowland, Mollick, Lifshitz, Kellogg, Rajendran, Krayer, Candelon, and Lakhani** Navigating the jagged technological frontier: Field experimental evidence of the effects of artificial intelligence on knowledge worker productivity and quality. *Organization Science*, 37(2).
- [5] **Brynjolfsson, Eric and Danielle Li and Lindsay Raymond** Generative AI at work. *The Quarterly Journal of Economics*, 140(2):889–942.
- [6] **Autor, David H, Frank Levy, and Richard J Murnane**, “The skill content of recent technological change: An empirical exploration,” *The Quarterly Journal of Economics*, 2003, 118 (4), 1279–1333.
- [7] **Acemoglu, Daron and David Autor**, “Skills, tasks and technologies: Implications for employment and earnings,” in *Handbook of Labor Economics*, Vol. 4, Elsevier, Amsterdam, 2011, pp. 1043–1171.
- [8] **Deming, David J and Kadeem Noray**, “Earnings dynamics, changing job skills, and STEM careers,” *The Quarterly Journal of Economics*, 2020, 135 (4), 1965–2005.
- [9] **Li, Ling**, “Reskilling and upskilling the future-ready workforce for industry 4.0 and beyond,” *Information Systems Frontiers*, 2022, pp. 1–16.
- [10] **Eloundou, Tyna, Sam Manning, Pamela Mishkin, and Daniel Rock**, “GPTs are GPTs: Labor market impact potential of LLMs,” *Science*, 2024, 384 (6702), 1306–1308.
- [11] **Mollick, Ethan R and Lilach Mollick**, “New Modes of Learning Enabled by AI Chatbots: Three Methods and Assignments,” 12 2022. Available at SSRN: <https://ssrn.com/abstract=4300783> or <http://dx.doi.org/10.2139/ssrn.4300783>.
- [12] **Agarwal, Nikhil, Alex Moehring, Pranav Rajpurkar, and Tobias Salz**, “Combining human expertise with artificial intelligence: Experimental evidence from radiology,” Technical Report, National Bureau of Economic Research 2023.

- 735 [13] **Dell’Acqua, Fabrizio**, “Falling asleep at the wheel: Human/AI Collaboration
736 in a Field Experiment on HR Recruiters,” Technical Report, Working paper 2022.
- 737 [14] **Carlini, Nicholas**, “A GPT-4 Capability Forecasting Challenge,” 2023. Avail-
738 able at: <https://nicholas.carlini.com/writing/llm-forecast/>.
- 739 [15] **Wiles, Emma and John J Horton**, “Generative AI and labor market matching
740 efficiency,” *Available at SSRN 5187344*, 2025.
- 741 [16] **Otis, Nicholas, Rowan P Clarke, Solene Delecourt, David Holtz, and
742 Rembrand Koning**, “The uneven impact of generative AI on entrepreneurial
743 performance,” *Available at SSRN 4671369*, 2023.
- 744 [17] **Noy, Shakked and Whitney Zhang**, “Experimental evidence on the pro-
745 ductivity effects of generative artificial intelligence,” *Science*, 2023, 381 (6654),
746 187–192.
- 747 [18] **Wiles, Emma, Zanele Munyikwa, and John Horton**, “Algorithmic writing
748 assistance on jobseekers’ resumes increases hires,” *Management Science*, 2025.
- 749 [19] **Peng, S., E. Kalliamvakou, P. Cihon, and M. Demirer**, “The impact of
750 AI on developer productivity: Evidence from GitHub Copilot,” arXiv preprint
751 arXiv:2302.06590, 2023.
- 752 [20] **Anderson, John R.**, ed., *Cognitive Skills and Their Acquisition*, Lawrence
753 Erlbaum Associates, Hillsdale, NJ, 1981.
- 754 [21] **DeKeyser, Robert**, “Skill acquisition theory,” in *Theories in Second Language
755 Acquisition*, Routledge, New York and London, 2020, pp. 83–104.
- 756 [22] **Brown, Peter C, Henry L Roediger III, and Mark A McDaniel**, *Make It
757 Stick: The Science of Successful Learning*, Harvard University Press, Cambridge,
758 2014.
- 759 [23] **Bjork, Elizabeth L and Robert A Bjork**, “Making things hard on yourself,
760 but in a good way: Creating desirable difficulties to enhance learning,” *Psychology
761 and the Real World: Essays Illustrating Fundamental Contributions to Society*,
762 2011, 2 (56–64).
- 763 [24] **Agarwal, R. and J. Prasad**, “A conceptual and operational definition of
764 personal innovativeness in the domain of information technology,” *Information
765 Systems Research*, 1998, 9 (2), 204–215.
- 766 [25] **Miron, E., M. Erez, and E. Naveh**, “Do personal characteristics and cultural
767 values that promote innovation, quality, and efficiency compete or complement
768 each other?,” *Journal of Organizational Behavior*, 2004, 25 (2), 175–199.
- 769 [26] **Jha, S. and S. S. Bhattacharyya**, “Learning orientation and performance

- orientation: Scale development and its relationship with performance,” *Global Business Review*, 2013, *14* (1), 43–54.
- [27] **Lee, David S**, “Training, Wages, and Sample Selection: Estimating Sharp Bounds on Treatment Effects,” *Review of Economic Studies*, 2009, *76* (3), 1071–1102.